

Project: Wrangling and Analyze Data

Data Gathering

In the cell below, gather **all** three pieces of data for this project and load them in the notebook. **Note:** the methods required to gather each data are different.

```
In [832... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import requests
import tweepy
from tweepy import OAuthHandler
import json
from timeit import default_timer as timer
import seaborn as sns
```

1. Directly download the WeRateDogs Twitter archive data (twitter_archive_enhanced.csv)

```
In [833... df_archive=pd.read_csv('twitter-archive-enhanced.csv')
df_archive.head(10)
```

| Out[833]: | tweet_id | in_reply_to_status_id | in_reply_to_user_id | timestamp | source |
|-----------|--------------------|-----------------------|---------------------|---------------------------|---|
| 0 | 892420643555336193 | NaN | NaN | 2017-08-01 16:23:56 +0000 | href="http://twitter.com/download/iphon |
| 1 | 892177421306343426 | NaN | NaN | 2017-08-01 00:17:27 +0000 | href="http://twitter.com/download/iphon |
| 2 | 891815181378084864 | NaN | NaN | 2017-07-31 00:18:03 +0000 | href="http://twitter.com/download/iphon |
| 3 | 891689557279858688 | NaN | NaN | 2017-07-30 15:58:51 +0000 | href="http://twitter.com/download/iphon |
| 4 | 891327558926688256 | NaN | NaN | 2017-07-29 16:00:24 +0000 | href="http://twitter.com/download/iphon |
| 5 | 891087950875897856 | NaN | NaN | 2017-07-29 00:08:17 +0000 | href="http://twitter.com/download/iphon |

| | | | | | |
|---|--------------------|-----|-----|---------------------------|---|
| 6 | 890971913173991426 | NaN | NaN | 2017-07-28 16:27:12 +0000 | href="http://twitter.com/download/iphon |
|---|--------------------|-----|-----|---------------------------|---|

| | | | | | |
|---|--------------------|-----|-----|---------------------------|---|
| 7 | 890729181411237888 | NaN | NaN | 2017-07-28 00:22:40 +0000 | href="http://twitter.com/download/iphon |
|---|--------------------|-----|-----|---------------------------|---|

| | | | | | |
|---|--------------------|-----|-----|---------------------------|---|
| 8 | 890609185150312448 | NaN | NaN | 2017-07-27 16:25:51 +0000 | href="http://twitter.com/download/iphon |
|---|--------------------|-----|-----|---------------------------|---|

| | | | | | |
|---|--------------------|-----|-----|---------------------------|---|
| 9 | 890240255349198849 | NaN | NaN | 2017-07-26 15:59:51 +0000 | href="http://twitter.com/download/iphon |
|---|--------------------|-----|-----|---------------------------|---|

1. Use the Requests library to download the tweet image prediction (image_predictions.tsv)

```
In [834... #url link
url='https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-predictions
response=requests.get(url)

file_name = "image-predictions.tsv"
with open(file_name, 'wb') as file:
    file.write(response.content)

#test file
df_image=pd.read_csv('image-predictions.tsv', sep='\t')
df_image.head(10)
len(df_image)
```

Out[834]: 2075

1. Use the Tweepy library to query additional data via the Twitter API (tweet_json.txt)

```
In [835... #pip install tweepy

consumer_key = ''
consumer_secret = ''
access_token = ''
access_secret = ''

auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth, wait_on_rate_limit=True)
```

```
In [836... # Query Twitter's API for JSON data for each tweet ID in the Twitter archive
#count = 0
#fails_dict = {}
#start = timer()
```

```
# Save each tweet's returned JSON as a new line in a .txt file
#with open('tweet_json.txt', 'w') as outfile:
#    # This loop will likely take 20-30 minutes to run because of Twitter's rate limit
#    for tweet_id in tweet_ids:
#        count += 1
#        print(str(count) + ": " + str(tweet_id))
#        try:
#            tweet = api.get_status(tweet_id, tweet_mode='extended',wait_on_rate_limit=T
#            print("Success")
#            json.dump(tweet._json, outfile)
#            outfile.write('\n')
#        except tweepy.TweepError as e:
#            print("Fail")
#            fails_dict[tweet_id] = e
#            pass
#end = timer()
#print(end - start)
#print(fails_dict)
#
```

In []:

In [837]...

```
df_count = []

with open('tweet-json.txt', 'r') as file:
    for row in file:
        tweet = json.loads(row)
        tweet_id = tweet['id']
        retweet_count = tweet['retweet_count']
        fav_count = tweet['favorite_count']
        df_count.append({
            'tweet_id':tweet_id,
            'retweet_count': retweet_count,
            'favorite_count': fav_count,})

df_count= pd.DataFrame(df_count)
df_count.head()
```

Out[837]:

| | tweet_id | retweet_count | favorite_count |
|---|--------------------|---------------|----------------|
| 0 | 892420643555336193 | 8853 | 39467 |
| 1 | 892177421306343426 | 6514 | 33819 |
| 2 | 891815181378084864 | 4328 | 25461 |
| 3 | 891689557279858688 | 8964 | 42908 |
| 4 | 891327558926688256 | 9774 | 41048 |

Assessing Data

In this section, detect and document at least **eight (8) quality issues and two (2) tidiness issue**. You must use **both** visual assessment programmatic assessement to assess the data.

Note: pay attention to the following key points when you access the data.

- You only want original ratings (no retweets) that have images. Though there are 5000+ tweets in the dataset, not all are dog ratings and some are retweets.
- Assessing and cleaning the entire dataset completely would require a lot of time, and is not necessary to practice and demonstrate your skills in data wrangling. Therefore, the requirements of this project

dtypes: float64(4), int64(3), object(10)
memory usage: 297.7+ KB

In [840... *#A visual assessment of the table*
df_archive_no_retweets.sample(5)

Out[840]:

| | tweet_id | in_reply_to_status_id | in_reply_to_user_id | timestamp | s |
|--|----------|-----------------------|---------------------|-----------|---|
| | 729 | 781955203444699136 | NaN | NaN | 2016-09-30 20:33:43 +0000 href="http://twitter.com/download/ip |
| | 1327 | 705975130514706432 | NaN | NaN | 2016-03-05 04:36:02 +0000 href="http://twitter.com/download/ip |
| | 631 | 794205286408003585 | NaN | NaN | 2016-11-03 15:51:10 +0000 href="http://twitter.com/download/ip |
| | 2250 | 667832474953625600 | NaN | NaN | 2015-11-20 22:30:44 +0000 href="http://twitter.com/download/ip |
| | 1290 | 708130923141795840 | NaN | NaN | 2016-03-11 03:22:23 +0000 href="http://twitter.com/download/ip |

In [841... *#Check duplicated for the Tweet id column*
df_archive_no_retweets.tweet_id.duplicated()

Out[841]:

| | |
|------|-------|
| 0 | False |
| 1 | False |
| 2 | False |
| 3 | False |
| 4 | False |
| | ... |
| 2351 | False |
| 2352 | False |
| 2353 | False |
| 2354 | False |
| 2355 | False |

Name: tweet_id, Length: 2117, dtype: bool

In [842... *#Have a look at Unique Values. And the extent of repetition and if there are illogical v*
df_archive_no_retweets.name.value_counts()

Out[842]:

| | |
|---------|-----|
| None | 622 |
| a | 55 |
| Lucy | 11 |
| Charlie | 11 |
| Cooper | 10 |
| | ... |
| Wishes | 1 |
| Rose | 1 |
| Theo | 1 |
| Fido | 1 |

Christopher 1
Name: name, Length: 956, dtype: int64

In [843... *#Have a look at the image table structure*
df_image.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   tweet_id    2075 non-null   int64
1   jpg_url     2075 non-null   object
2   img_num     2075 non-null   int64
3   p1          2075 non-null   object
4   p1_conf     2075 non-null   float64
5   p1_dog      2075 non-null   bool
6   p2          2075 non-null   object
7   p2_conf     2075 non-null   float64
8   p2_dog      2075 non-null   bool
9   p3          2075 non-null   object
10  p3_conf     2075 non-null   float64
11  p3_dog      2075 non-null   bool
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

In [844... *#Visual evaluation of the table of pictures*
df_image.sample(10)

Out[844]:

| | tweet_id | jpg_url | img_num | p1 | p1_conf |
|------|--------------------|---|---------|------------------|---------|
| 897 | 699775878809702401 | https://pbs.twimg.com/media/CbYac83W4AAUH1O.jpg | 1 | Dandie_Dinmont | 0.27168 |
| 131 | 668297328638447616 | https://pbs.twimg.com/media/CUZE4IWW4AAZmDf.jpg | 1 | king_penguin | 0.60674 |
| 1931 | 859074603037188101 | https://pbs.twimg.com/media/C-wLyufW0AA546l.jpg | 1 | revolver | 0.19029 |
| 1034 | 711732680602345472 | https://pbs.twimg.com/media/CeCVGEbUYAASeY4.jpg | 3 | dingo | 0.36687 |
| 768 | 689280876073582592 | https://pbs.twimg.com/media/CZDRTAPUoAEaqxF.jpg | 3 | Chihuahua | 0.63754 |
| 1039 | 712085617388212225 | https://pbs.twimg.com/media/CeHWFksXIAAypp.jpg | 2 | Shih-Tzu | 0.62512 |
| 922 | 702217446468493312 | https://pbs.twimg.com/media/Cb7HCMkWEAAV9zY.jpg | 1 | golden_retriever | 0.24241 |
| 1333 | 757729163776290825 | https://pbs.twimg.com/media/CWyD2HGUYAQ1Xa7.jpg | 2 | cash_machine | 0.80233 |
| 438 | 674422304705744896 | https://pbs.twimg.com/media/CVwHgbIWcAACWOD.jpg | 1 | golden_retriever | 0.96449 |
| 1252 | 747963614829678593 | https://pbs.twimg.com/media/CmFM7ngXEAEitfh.jpg | 1 | kelpie | 0.30767 |

In [845... *#chack that there are no duplicates*
df_image.tweet_id.duplicated()

Out[845]:

| | |
|------|-------|
| 0 | False |
| 1 | False |
| 2 | False |
| 3 | False |
| 4 | False |
| ... | |
| 2070 | False |
| 2071 | False |
| 2072 | False |
| 2073 | False |
| 2074 | False |

Name: tweet_id, Length: 2075, dtype: bool

In [846... *# Check the stats count table*

```
df_count.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2354 entries, 0 to 2353
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tweet_id              2354 non-null   int64
1   retweet_count         2354 non-null   int64
2   favorite_count        2354 non-null   int64
dtypes: int64(3)
memory usage: 55.3 KB
```

In [847... *#visual assessment of the sample*
df_count.sample(5)

Out[847]:

| | tweet_id | retweet_count | favorite_count |
|------|--------------------|---------------|----------------|
| 110 | 871032628920680449 | 3999 | 23255 |
| 975 | 750011400160841729 | 1035 | 3568 |
| 820 | 770093767776997377 | 3520 | 0 |
| 1254 | 710588934686908417 | 2107 | 4968 |
| 649 | 792883833364439040 | 4964 | 12666 |

In [848... *#Verify.. that there is no duplicated*
df_count.tweet_id.duplicated()

Out[848]:

| | |
|------|-------|
| 0 | False |
| 1 | False |
| 2 | False |
| 3 | False |
| 4 | False |
| | ... |
| 2349 | False |
| 2350 | False |
| 2351 | False |
| 2352 | False |
| 2353 | False |

Name: tweet_id, Length: 2354, dtype: bool

In [849... *#For further verification, the two tables have been merged to find out the missing value*
df_join = pd.merge(df_archive_no_retweets, df_image, on="tweet_id", how="left")

In [850... *#Look at the table after the merge*
df_join.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2117 entries, 0 to 2116
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_id                             2117 non-null   int64
1   in_reply_to_status_id                23 non-null     float64
2   in_reply_to_user_id                  23 non-null     float64
3   timestamp                            2117 non-null   object
4   source                              2117 non-null   object
5   text                                 2117 non-null   object
6   retweeted_status_id                  0 non-null      float64
7   retweeted_status_user_id            0 non-null      float64
8   retweeted_status_timestamp          0 non-null      object
9   expanded_urls                        2117 non-null   object
```

```

10 rating_numerator      2117 non-null    int64
11 rating_denominator    2117 non-null    int64
12 name                  2117 non-null    object
13 doggo                 2117 non-null    object
14 floofer               2117 non-null    object
15 pupper               2117 non-null    object
16 puppo                2117 non-null    object
17 jpg_url               1994 non-null    object
18 img_num               1994 non-null    float64
19 p1                    1994 non-null    object
20 p1_conf               1994 non-null    float64
21 p1_dog                1994 non-null    object
22 p2                    1994 non-null    object
23 p2_conf               1994 non-null    float64
24 p2_dog                1994 non-null    object
25 p3                    1994 non-null    object
26 p3_conf               1994 non-null    float64
27 p3_dog                1994 non-null    object

```

dtypes: float64(8), int64(3), object(17)

memory usage: 479.6+ KB

```

In [851... # Take a look at the table after excluding tweets that do not have an image
df_join(((
    df_join['jpg_url'].notnull()
))
)].info()

```

<class 'pandas.core.frame.DataFrame'>

Int64Index: 1994 entries, 0 to 2116

Data columns (total 28 columns):

| # | Column | Non-Null Count | Dtype |
|----|----------------------------|----------------|---------|
| 0 | tweet_id | 1994 non-null | int64 |
| 1 | in_reply_to_status_id | 23 non-null | float64 |
| 2 | in_reply_to_user_id | 23 non-null | float64 |
| 3 | timestamp | 1994 non-null | object |
| 4 | source | 1994 non-null | object |
| 5 | text | 1994 non-null | object |
| 6 | retweeted_status_id | 0 non-null | float64 |
| 7 | retweeted_status_user_id | 0 non-null | float64 |
| 8 | retweeted_status_timestamp | 0 non-null | object |
| 9 | expanded_urls | 1994 non-null | object |
| 10 | rating_numerator | 1994 non-null | int64 |
| 11 | rating_denominator | 1994 non-null | int64 |
| 12 | name | 1994 non-null | object |
| 13 | doggo | 1994 non-null | object |
| 14 | floofer | 1994 non-null | object |
| 15 | pupper | 1994 non-null | object |
| 16 | puppo | 1994 non-null | object |
| 17 | jpg_url | 1994 non-null | object |
| 18 | img_num | 1994 non-null | float64 |
| 19 | p1 | 1994 non-null | object |
| 20 | p1_conf | 1994 non-null | float64 |
| 21 | p1_dog | 1994 non-null | object |
| 22 | p2 | 1994 non-null | object |
| 23 | p2_conf | 1994 non-null | float64 |
| 24 | p2_dog | 1994 non-null | object |
| 25 | p3 | 1994 non-null | object |
| 26 | p3_conf | 1994 non-null | float64 |
| 27 | p3_dog | 1994 non-null | object |

dtypes: float64(8), int64(3), object(17)

memory usage: 451.8+ KB

```

In [852... #Random sample
df_join.sample(5)

```


Out[852]:

| | tweet_id | in_reply_to_status_id | in_reply_to_user_id | timestamp | s |
|------|--------------------|-----------------------|---------------------|---------------------------|--------------------------------------|
| 1600 | 676263575653122048 | NaN | NaN | 2015-12-14 04:52:55 +0000 | href="http://twitter.com/download/ip |
| 202 | 845306882940190720 | NaN | NaN | 2017-03-24 16:10:40 +0000 | href="http://twitter.com/download/ip |
| 1143 | 703382836347330562 | NaN | NaN | 2016-02-27 00:55:11 +0000 | href="http://twitter.com/download/ip |
| 1841 | 670840546554966016 | NaN | NaN | 2015-11-29 05:43:44 +0000 | href="http://twitter.com/download/ip |
| 1687 | 674269164442398721 | NaN | NaN | 2015-12-08 16:47:50 +0000 | href="http://twitter.com/download/ip |

5 rows × 28 columns

| | | |
|------------|--|------|
| In [853... | #Check unique values | |
| | df_join['rating_denominator'].value_counts() | |
| Out[853]: | 10 | 2099 |
| | 50 | 3 |
| | 11 | 2 |
| | 80 | 2 |
| | 70 | 1 |
| | 7 | 1 |
| | 150 | 1 |
| | 170 | 1 |
| | 20 | 1 |
| | 90 | 1 |
| | 40 | 1 |
| | 130 | 1 |
| | 110 | 1 |
| | 120 | 1 |
| | 2 | 1 |
| | Name: rating_denominator, dtype: int64 | |

| | | |
|------------|---|---|
| In [854... | #Visually check the column for the number of images, as well as the link for the images | |
| | df_join[["img_num","jpg_url"]].sort_values(by=['img_num'], ascending=False).head(10) | |
| Out[854]: | img_num | jpg_url |
| | 304 | 4.0 https://pbs.twimg.com/media/C3rN-lcWEAA9CmR.jpg |
| | 504 | 4.0 https://pbs.twimg.com/media/CvukbEkWAAAV-69.jpg |
| | 608 | 4.0 https://pbs.twimg.com/media/Cr2_6R8WAAAUMtc.jpg |
| | 283 | 4.0 https://pbs.twimg.com/media/C4UZLZLWYAA0dcs.jpg |

| | | |
|------|-----|--|
| 278 | 4.0 | https://pbs.twimg.com/media/C41st0bXAAE6MP8.jpg |
| 33 | 4.0 | https://pbs.twimg.com/media/DEi_N9qXYAAgEEw.jpg |
| 1315 | 4.0 | https://pbs.twimg.com/media/CZMJYCRVAAE35Wk.jpg |
| 273 | 4.0 | https://pbs.twimg.com/media/C4uLLGuUoAAkIHm.jpg |
| 181 | 4.0 | https://pbs.twimg.com/media/C8m3-iQVoAAETnF.jpg |
| 162 | 4.0 | https://pbs.twimg.com/media/C9px7jyVwAAAnmwN.jpg |

Quality issues

df_archive

1. We actually have 2117 tweets but the number of tweets that have df_image is 1994 - we have missing values
2. timestamp column Object type has to be modified and tweet_id column int64 to Object
3. In the Columns of the (name , doggo , flooter, pupper ,puppo), there are values with the name None, which means that they are empty values. The value must be modified and made empty
4. (retweeted_status_id , retweeted_status_user_id , retweeted_status_timestamp) These columns contain null values. They show retweeted answers being dropped
5. (in_reply_to_status_id ,in_reply_to_user_id) Columns contain missing values and do not have and cannot be used. They must be dropped
6. (rating_denominator) column contains a number of 2099 as a result of a value of 10, which is illogical and makes this information useless.
7. (source) column contains redundant and useless data that should be dropped

df_image

1. jpg_url column contains only one picture. We show in a column img_num that contains more than one picture

Tidiness issues

1. The table (df_archive , df_image , df_join) must be merged into one table
2. Merge the data from (doggo, puppo, pupper, floofer) into a single column. NEW

Cleaning Data

In this section, clean **all** of the issues you documented while assessing.

Note: Make a copy of the original data before cleaning. Cleaning includes merging individual pieces of data according to the rules of [tidy data](#). The result should be a high-quality and tidy master pandas DataFrame (or DataFrames, if appropriate).

```

#scope of work
#You only want original ratings (no retweets) that have images
#df_archive_new = df_archive[(df_archive['retweeted_status_id'].isnull() & df_archive['e

#copy data
df_archive_new = df_archive.copy()

# drop retweeted row
drop_row = df_archive_new[(df_archive_new['retweeted_status_id'].notnull())].index
df_archive_new.drop(drop_row, inplace=True)

# drop rows without images
drop_row = df_archive_new[(df_archive_new['expanded_urls'].isnull())].index
df_archive_new.drop(drop_row, inplace=True)

#copy image
df_image_new = df_image.copy()

#copy Likes and retweets counts
df_count_new = df_count.copy()

```

```

In [856.. #chack row retweeted
df_archive_new['tweet_id'][(df_archive_new['retweeted_status_id'].notnull())].count()

```

Out[856]: 0

```

In [857.. #chack row without images
df_archive_new['tweet_id'][(df_archive_new['expanded_urls'].isnull())].count()

```

Out[857]: 0

```

In [858.. #info data
df_archive_new.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2117 entries, 0 to 2355
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tweet_id                             2117 non-null   int64
1   in_reply_to_status_id                 23 non-null     float64
2   in_reply_to_user_id                  23 non-null     float64
3   timestamp                             2117 non-null   object
4   source                               2117 non-null   object
5   text                                 2117 non-null   object
6   retweeted_status_id                  0 non-null      float64
7   retweeted_status_user_id             0 non-null      float64
8   retweeted_status_timestamp           0 non-null      object
9   expanded_urls                         2117 non-null   object
10  rating_numerator                      2117 non-null   int64
11  rating_denominator                   2117 non-null   int64
12  name                                 2117 non-null   object
13  doggo                               2117 non-null   object
14  floofer                             2117 non-null   object
15  pupper                              2117 non-null   object
16  puppo                               2117 non-null   object
dtypes: float64(4), int64(3), object(10)
memory usage: 297.7+ KB

```

Issue #1:

Define

The table (df_archive , df_image , df_join) must be merged into one table

Code

```
In [859... #Merging tables into one table and linking them by Tweet ID
df_Tweet_stats2 = []
df_Tweet_stats2 = df_archive_new.merge(df_image_new , on="tweet_id" ,how="left")
df_Tweet_stats3 = df_Tweet_stats2.merge( df_count_new , on="tweet_id" ,how="left")
```

Test

```
In [860... #
df_Tweet_stats3.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2117 entries, 0 to 2116
Data columns (total 30 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   tweet_id              2117 non-null   int64
 1   in_reply_to_status_id  23 non-null     float64
 2   in_reply_to_user_id    23 non-null     float64
 3   timestamp              2117 non-null   object
 4   source                 2117 non-null   object
 5   text                   2117 non-null   object
 6   retweeted_status_id    0 non-null      float64
 7   retweeted_status_user_id 0 non-null      float64
 8   retweeted_status_timestamp 0 non-null      object
 9   expanded_urls          2117 non-null   object
10   rating_numerator        2117 non-null   int64
11   rating_denominator      2117 non-null   int64
12   name                   2117 non-null   object
13   doggo                  2117 non-null   object
14   floofer                2117 non-null   object
15   pupper                 2117 non-null   object
16   puppo                  2117 non-null   object
17   jpg_url                1994 non-null   object
18   img_num                 1994 non-null   float64
19   p1                      1994 non-null   object
20   p1_conf                 1994 non-null   float64
21   p1_dog                  1994 non-null   object
22   p2                      1994 non-null   object
23   p2_conf                 1994 non-null   float64
24   p2_dog                  1994 non-null   object
25   p3                      1994 non-null   object
26   p3_conf                 1994 non-null   float64
27   p3_dog                  1994 non-null   object
28   retweet_count           2117 non-null   int64
29   favorite_count          2117 non-null   int64
dtypes: float64(8), int64(5), object(17)
memory usage: 512.7+ KB
```

Issue #2:

Define

In the Columns of the (name , doggo , floofer , pupper ,puppo), there are values with the name None, which means that they are empty values. The value must be modified and made empty

```
In [861... #
Column_None = ["name" , "doggo" , "floofer", "pupper" , "puppo"]
```

```
df_Tweet_stats3[Column_None] = df_Tweet_stats3[Column_None].replace({ "None": np.nan })
```

Code

```
In [862... #
Column_None = ["name" , "doggo" , "floofer", "pupper" ,"puppo"]
df_Tweet_stats3[Column_None] = df_Tweet_stats3[Column_None].replace({ "None": np.nan })
```

Test

```
In [863... if ("None" in set(df_Tweet_stats3[Column_None])) :
    print("recodeing")
else :
    print('done')
```

done

Issue #3 :

Define

Merge the data from (doggo, puppo, pupper, floofer) into a single column (category_dog)

```
In [864... # columns
columns = ['doggo', 'floofer', 'pupper', 'puppo']

# Merge the data from (doggo, puppo, pupper, floofer) into a single column (category_dog)
df_Tweet_stats3['category_dog'] = df_Tweet_stats3[columns].apply(lambda x: ','.join(x.dr

#fix null value
df_Tweet_stats3['category_dog'] = df_Tweet_stats3['category_dog'].replace({ "": np.nan })
```

Test

```
In [865... if(df_Tweet_stats3['category_dog'].count()== (df_Tweet_stats3['doggo'].count()+df_Tweet_
    print("done ")
else :
    print(df_Tweet_stats3['category_dog'].value_counts())
```

```
pupper      222
doggo        72
puppo        23
doggo,pupper  10
floofer        9
doggo,puppo    1
doggo,floofer  1
Name: category_dog, dtype: int64
```

This is a quality problem, which is the presence of more than one value for the same line. We will solve it next (doggo,pupper ,doggo,puppo, doggo,floofer)

Issue #4:

Define

category_dog columns: This is a quality problem, which is the presence of more than one value for the same line(doggo,pupper doggo,puppo, doggo,floofer) we will drop row

Code

```
In [866.. #value drop
category_drop = ['doggo,pupper' , 'doggo,puppo' , 'doggo,floofer']

#
category_drop = df_Tweet_stats3[df_Tweet_stats3['category_dog'].isin(category_drop) == T

#
df_Tweet_stats3.drop(category_drop , inplace=True)
```

Test

```
In [867.. #value drop
category_drop = ['doggo,pupper' , 'doggo,puppo' , 'doggo,floofer']

if df_Tweet_stats3['category_dog'][df_Tweet_stats3['category_dog'].isin(category_drop)].
    #Drop columns 'doggo', 'floofer', 'pupper', 'puppo'
    df_Tweet_stats3 = df_Tweet_stats3.drop(columns, axis=1)
    print('done and Drop columns')

else :
    print("recodeing")
```

done and Drop columns

Issue #5:

Define

Drop the following columns: retweeted_status_id , retweeted_status_user_id , retweeted_status_timestamp , jpg_url , in_reply_to_status_id, in_reply_to_user_id ,source

Code

```
In [868.. #Drop columns
drop_list = [ 'retweeted_status_id' , 'retweeted_status_user_id' , 'retweeted_status_tim

df_Tweet_stats3 = df_Tweet_stats3.drop(drop_list, axis=1)
```

Test

```
In [869.. #
if set(drop_list).issubset(df_Tweet_stats3.columns):
    print('not working re codeing')
else :
    df_Tweet_stats3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2105 entries, 0 to 2116
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   tweet_id              2105 non-null   int64
1   timestamp              2105 non-null   object
2   text                   2105 non-null   object
3   expanded_urls          2105 non-null   object
```

```

4   rating_numerator      2105 non-null   int64
5   rating_denominator    2105 non-null   int64
6   name                  1490 non-null   object
7   img_num               1983 non-null   float64
8   p1                    1983 non-null   object
9   p1_conf               1983 non-null   float64
10  p1_dog                1983 non-null   object
11  p2                    1983 non-null   object
12  p2_conf               1983 non-null   float64
13  p2_dog                1983 non-null   object
14  p3                    1983 non-null   object
15  p3_conf               1983 non-null   float64
16  p3_dog                1983 non-null   object
17  retweet_count         2105 non-null   int64
18  favorite_count        2105 non-null   int64
19  category_dog          326 non-null   object
dtypes: float64(4), int64(5), object(11)
memory usage: 345.4+ KB

```

Issue #6:

Define

Column type timestamp to datetime and p1_dog,p2_dog,p3_dog to bool and tweet_id to object

Code

```

In [870... #
to_bool = [ 'p1_dog','p2_dog','p3_dog' ]
to_timestamp = 'timestamp'

df_Tweet_stats3[to_bool] = df_Tweet_stats3[to_bool].astype('bool')
df_Tweet_stats3[to_timestamp] = pd.to_datetime(df_Tweet_stats3[to_timestamp])
df_Tweet_stats3['tweet_id'] = df_Tweet_stats3['tweet_id'].astype('object')

```

Test

```

In [871... if df_Tweet_stats3[to_bool[0]].dtypes == 'bool' :
    print('done')
if df_Tweet_stats3[to_bool[1]].dtypes == 'bool' :
    print('done')
if df_Tweet_stats3[to_bool[2]].dtypes == 'bool' :
    print('done')
if df_Tweet_stats3[to_timestamp].dtypes == 'datetime64[ns, UTC]' :
    print('done')
if df_Tweet_stats3['tweet_id'].dtypes == 'object' :
    print('done')

```

```

done
done
done
done
done

```

Storing Data

Save gathered, assessed, and cleaned master dataset to a CSV file named "twitter_archive_master.csv".

```
In [872... df_Tweet_stats3.to_csv('Twitter_archive_master.csv')
Twitter_archive_master=pd.read_csv('Twitter_archive_master.csv')
```

Analyzing and Visualizing Data

In this section, analyze and visualize your wrangled data. You must produce at least **three (3) insights and one (1) visualization.**

```
In [873... ##
Twitter_archive_master.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2105 entries, 0 to 2104
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             2105 non-null   int64
1   tweet_id               2105 non-null   int64
2   timestamp              2105 non-null   object
3   text                   2105 non-null   object
4   expanded_urls          2105 non-null   object
5   rating_numerator       2105 non-null   int64
6   rating_denominator     2105 non-null   int64
7   name                   1490 non-null   object
8   img_num                1983 non-null   float64
9   p1                     1983 non-null   object
10  p1_conf                 1983 non-null   float64
11  p1_dog                  2105 non-null   bool
12  p2                      1983 non-null   object
13  p2_conf                 1983 non-null   float64
14  p2_dog                  2105 non-null   bool
15  p3                      1983 non-null   object
16  p3_conf                 1983 non-null   float64
17  p3_dog                  2105 non-null   bool
18  retweet_count           2105 non-null   int64
19  favorite_count          2105 non-null   int64
20  category_dog            326 non-null    object
dtypes: bool(3), float64(4), int64(6), object(8)
memory usage: 302.3+ KB
```

```
In [874... #After converting the dataframe into a file, you must re-assign the columns types again
to_bool = [ 'p1_dog','p2_dog','p3_dog' ]
to_timestamp = 'timestamp'

Twitter_archive_master[to_bool] = Twitter_archive_master[to_bool].astype('bool')
Twitter_archive_master[to_timestamp] = pd.to_datetime(Twitter_archive_master[to_timestamp])
Twitter_archive_master['tweet_id'] = Twitter_archive_master['tweet_id'].astype('object')
```

```
In [875... #
Twitter_archive_master.describe()
```

Out[875]:

| | Unnamed: 0 | rating_numerator | rating_denominator | img_num | p1_conf | p2_conf | p3_conf |
|-------|-------------|------------------|--------------------|-------------|-------------|--------------|--------------|
| count | 2105.000000 | 2105.000000 | 2105.000000 | 1983.000000 | 1983.000000 | 1.983000e+03 | 1.983000e+03 |
| mean | 1061.045131 | 12.251781 | 10.504038 | 1.203732 | 0.593112 | 1.346413e-01 | 6.034123e-02 |
| std | 611.422625 | 40.389622 | 7.125979 | 0.561861 | 0.271886 | 1.007247e-01 | 5.089379e-02 |
| min | 0.000000 | 0.000000 | 2.000000 | 1.000000 | 0.044333 | 1.011300e-08 | 1.740170e-10 |
| 25% | 532.000000 | 10.000000 | 10.000000 | 1.000000 | 0.362715 | 5.417505e-02 | 1.624950e-02 |

| | | | | | | | |
|------------|-------------|-------------|------------|----------|----------|--------------|--------------|
| 50% | 1064.000000 | 11.000000 | 10.000000 | 1.000000 | 0.587342 | 1.176080e-01 | 4.965610e-02 |
| 75% | 1590.000000 | 12.000000 | 10.000000 | 1.000000 | 0.843635 | 1.953115e-01 | 9.158675e-02 |
| max | 2116.000000 | 1776.000000 | 170.000000 | 4.000000 | 1.000000 | 4.880140e-01 | 2.734190e-01 |

Insights:


nsights 1:

Do you interact with tweets more by retweeting or favorite?

```
In [876... #sum of your favorites and retweets
Twitter_sum =Twitter_archive_master[['retweet_count','favorite_count']].sum()
Twitter_sum
```

```
Out[876]: retweet_count      5921833
favorite_count    18733074
dtype: int64
```

```
In [877... pie = Twitter_sum.plot(kind='pie',fontsize=20, legend=False, autopct='%.f', subplots=False)
```



It shows us that favorites is more than retweeting by a large margin

nsights 2:

What are the most 10 common dog breeds in machine learning?

```
In [878... dogs_type_best_p = []
dogs_type = Twitter_archive_master
dogs_type = dogs_type.reset_index()
for index, row in dogs_type.iterrows():
    if row['p1_dog'] == True :
        dogs_type_best_p.append({
            'tweet_id':row['tweet_id'],
            'dog_type': row['p1']})
    elif row['p2_dog'] == True :
        dogs_type_best_p.append({
            'tweet_id':row['tweet_id'],
            'dog_type': row['p2']})
    elif row['p3_dog'] == True :
        dogs_type_best_p.append({
            'tweet_id':row['tweet_id'],
            'dog_type': row['p3']})

dogs_type_best_p= pd.DataFrame(dogs_type_best_p)
```


```
In [879... dogs_type = dogs_type.merge( dogs_type_best_p , on="tweet_id" ,how="left")
```

```
In [880... hist_dog_hist = dogs_type.dog_type.value_counts()
```

```
In [881... hist_dog_hist= pd.DataFrame(hist_dog_hist)
```

```
In [882... hist_dog_hist.iloc[:10 ].plot(kind='bar', fontsize=15,figsize=(14,8),title='10 common do
```

```
Out[882]: <AxesSubplot:title={'center':'10 common dog breeds in machine learning'}>
```



nsights 3:

Are the tweets with videos more interactive than other tweets?

```
In [883... #Photo Tweets
twitter_photo = Twitter_archive_master[Twitter_archive_master['expanded_urls'].str.contains('#Video Tweets')
twitter_video = Twitter_archive_master[Twitter_archive_master['expanded_urls'].str.contains('#Photo Tweets')]
```

```
In [884... twitter_photo[['expanded_urls', 'name', 'retweet_count', 'favorite_count']].describe()
```

```
Out[884]:
```

| | retweet_count | favorite_count |
|-------|---------------|----------------|
| count | 2032.000000 | 2032.000000 |
| mean | 2581.029528 | 8483.309055 |
| std | 3678.925470 | 11032.480126 |
| min | 16.000000 | 81.000000 |
| 25% | 620.750000 | 1978.500000 |
| 50% | 1359.500000 | 4053.500000 |
| 75% | 3154.000000 | 11075.500000 |
| max | 48265.000000 | 132810.000000 |

```
In [885... twitter_video[['expanded_urls', 'name', 'retweet_count', 'favorite_count']].describe()
```

```
Out[885]:
```

| | retweet_count | favorite_count |
|--|---------------|----------------|
|--|---------------|----------------|

| | | |
|--------------|--------------|---------------|
| count | 73.000000 | 73.000000 |
| mean | 9276.452055 | 20479.315068 |
| std | 14632.284991 | 26988.841923 |
| min | 388.000000 | 934.000000 |
| 25% | 1224.000000 | 3047.000000 |
| 50% | 3285.000000 | 9555.000000 |
| 75% | 9907.000000 | 25057.000000 |
| max | 79515.000000 | 131075.000000 |

Through the previous results, it is clear that interaction with tweets that contain video is more than that of images

Visualization

```
In [886.. # mean photo retweet_count favorite_count
photo_mean = twitter_photo[['retweet_count', 'favorite_count']].mean()

# mean video retweet_count favorite_count
video_mean = twitter_video[['retweet_count', 'favorite_count']].mean()
names = ['retweet_count', 'favorite_count']

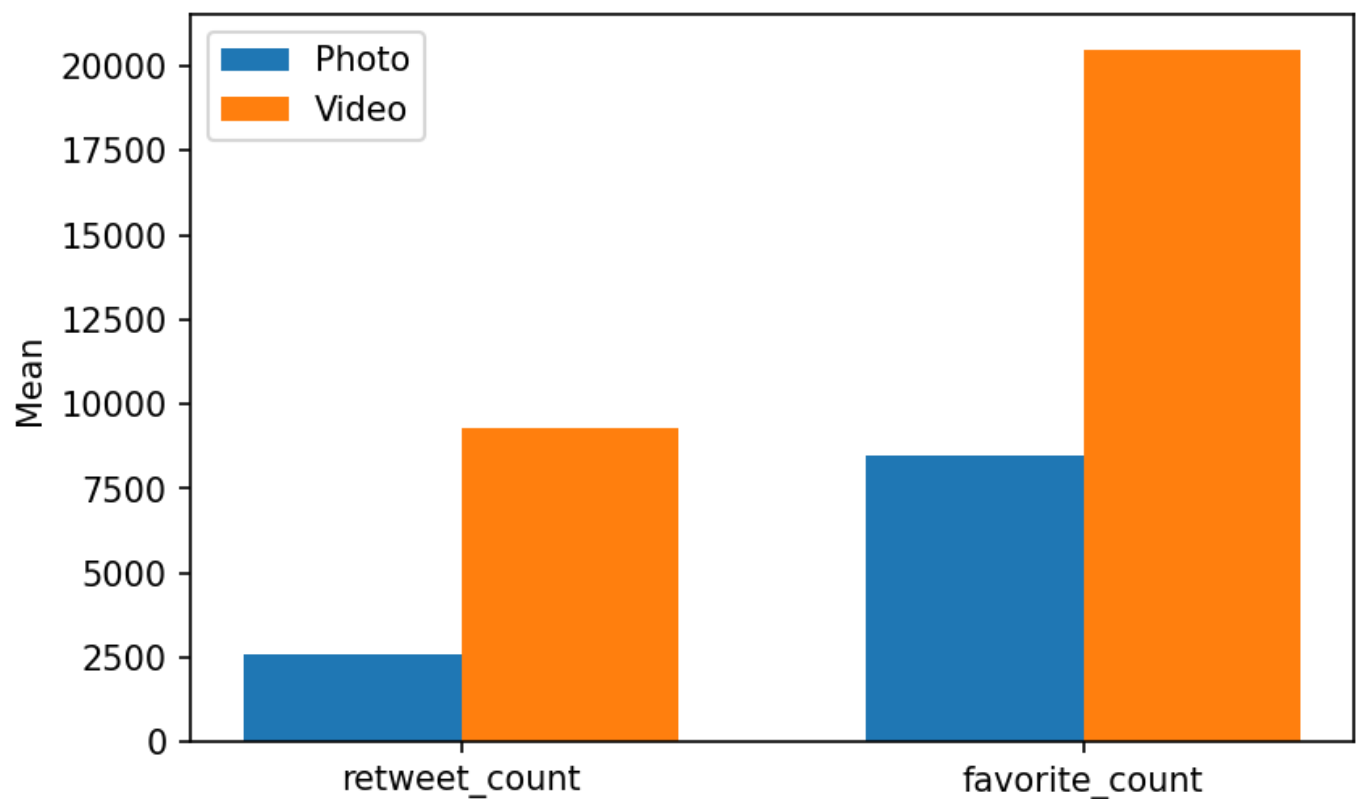
fig = plt.figure(figsize=(6,5), dpi=150)
left, bottom, width, height = 0.1, 0.3, 0.8, 0.6
ax = fig.add_axes([left, bottom, width, height])

width = 0.35
ticks = np.arange(len(names))
ax.bar(ticks, photo_mean, width, label='Photo')
ax.bar(ticks + width, video_mean, width, align="center", label='Video')

ax.set_ylabel('Mean')
ax.set_title('Photo Tweets vs Video Tweets')
ax.set_xticks(ticks + width/2)
ax.set_xticklabels(names)

ax.legend(loc='best')
plt.show()
```

Photo Tweets vs Video Tweets



```
In [887... # mean photo Rating Numerator Rating Denominator
photo_mean = twitter_photo[['rating_numerator','rating_denominator']].mean()

# mean video Rating Numerator Rating Denominator
video_mean = twitter_video[['rating_numerator','rating_denominator']].mean()
names = ['Rating Numerator','Rating Denominator']

fig = plt.figure(figsize=(6,5), dpi=150)
left, bottom, width, height = 0.1, 0.3, 0.8, 0.6
ax = fig.add_axes([left, bottom, width, height])

width = 0.35
ticks = np.arange(len(names))
ax.bar(ticks, photo_mean, width, label='Photo')
ax.bar(ticks + width, video_mean, width, align="center",label='Video')

ax.set_ylabel('Mean')
ax.set_title('Photo Tweets vs Video Tweets')
ax.set_xticks(ticks + width/2)
ax.set_xticklabels(names)

ax.legend(loc='best')
plt.show()
```

Photo Tweets vs Video Tweets

